

Commodities

category

prediction

Students

Nesterov Grigoriy

Aleksandr Vashchenko

Georgy Andryushchenko

Course

MLOps Engineering

Semester

Summer

Year

2024

Contents

1	Introduction	1
2	Business and Data Understanding	2
2.1	Terminology	2
2.2	Scope of the ML project	4
2.3	Success criteria	5
2.4	Data collection	6
2.5	Data quality verification	6
2.6	Project feasibility	11
2.7	Project plan	14

1. Introduction

In recent years, marketplaces have taken a significant place in the Russian market. COVID-19 has proven to be one of the most significant drivers of interest growth in such services. In 2023 alone, total expenditures on marketplaces increased by 50%. Niche marketplaces are also showing substantial growth[1].

However, marketplaces are built not only around customers, sellers are equally important. Therefore, it becomes crucial to help sellers choose the most appropriate category for the goods they are trying to sell. This way, customers are more likely to find the product in the category they expect it to be, and sellers do not need to spend a lot of time selecting the best category out of hundreds available.

To address this challenge, our project utilizes the «Online Retail Dataset» from Kaggle. This dataset contains information about purchases in an online store over a certain period, allowing us to conduct an in-depth analysis and identify key patterns that can help improve the category selection process for products.

2. Business and Data Understanding

In this project, the primary business problem is to increase seller loyalty by providing them with product category recommendations based on product features, including its description. It is essential to understand that proper product categorization allows customers to find the desired products more quickly, which increases the likelihood of purchase, and consequently, seller satisfaction and their willingness to continue working with the platform. The current situation shows that sellers often face difficulties in selecting the appropriate category for their products, leading to reduced sales and dissatisfaction. Therefore, we need to answer the question: "Is it possible to predict the product category with 90% accuracy based on a product description?". Using data from the "Online Retail Dataset" on the Kaggle platform will enable us to analyze and identify key parameters for building a product category prediction model [2].

2.1 Terminology

2.1.1 Business terminology

- Product - item offered for a sale on a marketplace;
- Seller - person or organization gaining profit by products sales;
- Product feature - textual, numerical or visual representation of product property;
- Product description - detailed description of the product qualities, design purposes, and features, written by the seller;
- Product card - organized collection of product features containing all sufficient product features which are essential for customers and are mandatory for product publication on a marketplace;
- Product category - product feature that is used by marketplace to organize the catalogue of products;
- Filtering system - set of manually designed and updated rules ensuring the validity of product features and descriptions, their alignment with the values of the marketplace rules, governmental laws, and social virtues.

2.1.2 ML terminology

- Machine learning model - a software program capable of making predictions within a specific domain, such as predicting product sales by category in online retail. Example: A machine learning model trained on historical customer purchase data can predict future sales volumes for different product categories based on factors like seasonality, customer demographics, and marketing promotions.
- Dashboard - an interactive website or application designed to display and visualize key properties and insights derived from relevant data. For example, accuracy of ML model by predict sales in stores. Example: An e-commerce analytics dashboard showing metrics such as daily sales trends, top-selling products, customer

demographics, and profitability by product category. Example: In an online retail platform, the confusion matrix helps to visualize how well a sales prediction model identifies successful product category sales (positive) and unsuccessful sales (negative) based on actual vs. predicted outcomes.

- Confusion matrix - the confusion matrix is a table used to describe the performance of a classification model. It contains four entries:
 - TP (True Positive): The number of correctly predicted positive samples.
 - TN (True Negative): The number of correctly predicted negative samples.
 - FP (False Positive): The number of incorrectly predicted positive samples.
 - FN (False Negative): The number of incorrectly predicted negative samples.

	Predicted Positive	Predicted Negative
Actual Positive	TP	FN
Actual Negative	FP	TN

Table 1: Confusion Matrix

- Precision - the ratio of correctly predicted positive observations to the total predicted positives. It is calculated using the formula: $\text{Precision} = \frac{TP}{TP+FP}$. Example: In an e-commerce platform, precision measures how accurately the model predicts successful sales (positive predictions) for specific product categories among all predicted sales.
- Recall - the ratio of correctly predicted positive observations to all observations in the actual class (also known as Sensitivity or True Positive Rate). It is calculated using the formula: $\text{Recall} = \frac{TP}{TP+FN}$. Example: In an online retail recommendation system, recall measures how well the model captures all successful sales (true positives) for specific product categories among all actual sales.
- F1-Score - the harmonic mean of Precision and Recall. It is used to evaluate the performance of a classification model, especially when the balance between Precision and Recall is important. The formula for the F1-score is defined as follows: $F1 = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$. Example: In an e-commerce platform, the F1-score helps assess how well the model balances the ability to accurately predict successful product category sales (precision) with capturing all actual successful sales (recall).
- Accuracy - a metric that measures the overall performance of a classification model. It calculates the ratio of correctly predicted observations to the total number of observations. The formula for accuracy is defined as follows: $\text{Accuracy} = \frac{TP+TN}{TP+TN+FP+FN}$. Example: In an online retail sales forecasting system, accuracy quantifies how well the model correctly predicts both successful (sales) and unsuccessful (non-sales) outcomes for product categories.
- Weighted accuracy - a metric that considers the contribution of each class (category) to the overall accuracy of the model, weighted by the size of each class. The

formula for calculating weighted accuracy is expressed as follows: $\text{Accuracy}_{\text{weighted}} = \frac{\sum_{i=1}^N w_i \cdot \text{Accuracy}_i}{\sum_{i=1}^N w_i}$ where:

- N - is the total number of classes (categories).
- w_i - is the weight assigned to class i , typically proportional to the class distribution in the dataset.
- Accuracy_i - is the accuracy for class i , calculated as $\frac{TP_i + TN_i}{TP_i + TN_i + FP_i + FN_i}$, where TP_i , TN_i , FP_i , and FN_i are the counts for true positives, true negatives, false positives, and false negatives, respectively, for class i .

Example: In an online retail platform with multiple product categories, weighted accuracy evaluates how accurately the sales prediction model forecasts sales for each category, giving more weight to categories with higher sales volumes or strategic importance.

2.2 Scope of the ML project

2.2.1 Background

Ozon is one of the most famous and largest marketplaces in Russia. Based on their official tutorial [3], the seller must fill in the following entries while registering a new product:

1. Category;
2. Product title;
3. Stock keeping unit;
4. Price, fees, and discounts;
5. Packaging parameters - size, weights, variants of packages (number of items in one package; length of the rope, cable, hose, etc);
6. Photos;
7. Expiration dates;
8. Brand;
9. Color;
10. And few other dozens of product features.

As we could see, there are many entries to be filled in and while some of them are fairly simple to fill in - like packaging parameters, others require in-depth understanding of marketplace tag systems - including categories.

The major problem with categories is caused by the weak borders between the categories and subcategories - the same pair of shoes might be categorized both as a sports shoe and casual shoes. Therefore, we as a marketplace platform may simplify the category selection by suggesting the most appropriate category based on other essential features of the product.

We believe that with the use of the simplified category selection approach sellers are less likely to put their product in an inappropriate category and thus get higher sells. Therefore, our platform fees will also increase.

2.2.2 Business problem

Sellers set up an incorrect product category to result in an overall loss of profit for both the platform and the sellers.

2.2.3 Business objectives

Taking into account the analysis above we may formulate the business objective as follows: increase the loyalty of the sellers by providing them with product category recommendation based on product features, including its description. Therefore, we need to answer the question: "Is it possible to predict the product category with 90% accuracy based on a product description?".

2.2.4 ML objectives

Create a classical ML model that uses basic NLP to predict the product category based on its textual description, product name, available variations, and price. Ideally, preserve model interpretability. Avoid use of large models to keep deployment, operation, and maintenance cheap.

2.3 Success criteria

2.3.1 Business Success Criteria

- Simplify the product card creation via suggestion of the most likely product category. Representative sellers group can evaluate the final service convenience;
- Reduce new sellers churn by at least 1%.

2.3.2 ML Success Criteria

- Achieve at least 90% accuracy in category prediction on unseen data;
- Use only classical statistical and ML models to ensure low computational demands;
- Use only classical NLP techniques - such as TF-IDF and Word2Vec embeddings to keep models simple and interpretable.

2.3.3 Economic Success Criteria

- Increase in Revenue: Achieve a 5% increase in total platform revenue attributed to enhanced product categorization accuracy and efficiency.
- Cost Savings: Reduce operational costs related to manual categorization efforts by 15% through automation and improved accuracy.

- Seller Retention Improvement: Improve seller retention rates by at least 3% due to better product visibility and placement within appropriate categories.

2.4 Data collection

2.4.1 Data collection report

Kaggle[2] dataset was selected as the data source. The data were provided in a csv file in tabular form. The data contain the following structure: 12 columns, 464433 rows. To encourage reproducibility, we also saved a copy of the data on our [Yandex disk](#).

2.4.2 Data version control report

To ensure the reliability and reproducibility of our machine learning models, we use Data Version Control (DVC). This tool helps us manage large datasets, track data versions, and log changes. Each version of the data, such as the initial raw data, cleaned data, and feature-engineered data, is tagged and documented. Changes to the data are recorded with detailed descriptions.

We back up data using a combination of cloud storage and on-premises servers, ensuring redundancy and accessibility. DVC stores metadata in our Git repository, while actual data is stored remotely. Historical data is retained for audit and compliance purposes. Access to the data is restricted to authorized team members, with roles and permissions managed to ensure data integrity.

The importance of data version control includes:

- Retraining models based on new data or approaches.
- Addressing model degradation over time.
- Quickly rolling back to previous versions if needed.
- Ensuring compliance with corporate or government regulations.

2.5 Data quality verification

2.5.1 Data description

In total, provided sample contains 464433 records distributed along 12 columns:

1. itemid - unique item identifier. It appears to be actually unique;
2. shopid - shop identifier. In total we have information from 7856 shops;
3. item_name - name of the product that is used during its sale. We found 118346 unique item names but chances are that some of them are just different names for the same item;
4. item_description - textual description of item provided by a seller;
5. item_variation - array of possible item variations - different colors, package sizes, shoe size and so on;

6. price - price of the item in USD;
7. stock - how much items are still available;
8. category - to which of 21 unique categories the item belongs;
9. cb_option - binary feature indicating whether the item can be sold cross-border. Only 11% of items can not be sold to other countries;
10. is_preferred - binary feature indicating whether the item is sold by a shop focusing on this category and having an appropriate license or working with a manufacturer directly. Only 4% of the items are sold by preferred shops;
11. sold_count - how many items the shop has sold so far;
12. item_creation_date - the date when the seller has created this item card. It ranges from Jan 1, 2016 to Sept 9, 2017.

2.5.2 Data exploration

First of all, we decided to see how many repeated description do we have in total - and we found out that 75.5% of descriptions are used at least twice. It means that different sellers are likely to use the same descriptions. This claim is strongly supported by the fact 17.6% of the item name - item description pairs are unique. Still, the number of unique descriptions is pretty high and thus we need to apply some NLP pipelines to efficiently utilize that data.

We further discover that on average, the item title contains 60 characters and is 10 times shorter than the item description (that contains on average 635 characters). Distribution of lengths are presented of Figures 1 and 2. We may notice that the distribution of name lengths rapidly drops after about 70 characters while the distribution of description lengths resembles a truncated normal distribution. From this we may conclude that descriptions are longer than titles but still remain quite shord.




Figure 1: Distribution of description lengths




Figure 2: Distribution of description lengths

We also decided to analyze the number of unique descriptions per item. The results are presented on a figure 3. Analyzing the plot we may conclude that majority of items have at most two unique descriptions.




Figure 3: Distribution of the number of unique descriptions for items

2.5.3 Data requirements

The data should be defined in the following order:

1. itemid - unique numeric item identifier between 1 and more;
2. shopid - shop numeric item identifier between 1 and more;
3. item_name - string data type between 5 and more length;
4. item_description - string data type between 10 and more length;
5. item_variation - array of possible item variations between 0 and more length;
6. price - float number between 0 and 9000000;
7. stock - numerical item between 0 and 400000;
8. category - string data type with length between 5 and more;
9. cb_option - binary feature 0 or 1;
10. is_preferred - binary feature 0 or 1;
11. sold_count - integer numeric between 0 and 8007;
12. item_creation_date - date type of data between Jan 1, 2016 and Sept 9, 2017.

2.5.4 Data quality verification report

Analyzing the quality of the data we came up with the following conclusions:

1. Chances are we have only top level categories available on a market place in this sample - usually market place also have a hierarchy of categories. For example, shoes might be subdivided by gender, by expected usage scenario (sports, official events, daily, and so on), season (summer, winter, autumn or spring), age group and in many other ways;
2. There are 1111 (about 2% of the data) rows having missing values represented by NaNs. These missing values appear in one columns - item_name, item_description, and category.

2.6 Project feasibility

2.6.1 Inventory of resources

List the resources available to the project including:

- Grigoriy Nesterov - Data Engineer;
- Georgy Andryushchenko - ML Engineer;
- Aleksandr Vashchenko - Data Scientist;
- Ozon Sellers educational platform[3];
- Extract of the data from online retail platform[2].

2.6.2 Requirements, assumptions and constraints

The key requirements of the project include:

1. On-time deployment - before 19th of July;
2. High quality of suggestions.

To guarantee on time deployment the time available (8 weeks) must be efficiently distributed. Table 2 contains expected project completion plan.

Week(-s)	What should be accomplished by the end of the week?
1	Business and Data Understanding
2-3	Data engineering/Preparation
3-4	Model engineering
5	Model validation
6	Model deployment and Model monitoring and maintenance

Table 2: Expected project execution plan

Some of the stages have completion time significantly overestimated to give some freedom for creativity and exploration, if necessary. Otherwise this time can be used for quality improvement of our service.

To ensure the high quality of suggestions we have to rely on several assumptions to be true:

1. Provided data sample provides a high quality representation of typical product cards;
2. Available features contain all the information necessary to properly predict the item category;
3. Distribution of categories in the data properly represents that of the real data.

Provided data sample is in public domain and therefore can be freely used without any litigation. Data is anonymized and does not include any personal information so even if the sample leaks nobody is harmed. All of the sudden, the volume of the data might be insufficient to use state-of-the-art natural language processing approaches involving neural networks.

Further, our project team is significantly limited in choice of tools as administration of IU Hadoop cluster is a responsibility of other authority.

2.6.3 Risks and contingencies

We organize major risks and their mitigation strategies in the Table [3](#).

Data induced risks	
Risk	Mitigation strategy
Non-representative data sample	Intensive work with customer to ensure high-quality of a new data sample
Outdated sample	Obtain a fresh sample or design a thorough model assessment and controlled deployment
Imbalanced data	Use re-sampling techniques
Model induced risks	
Risk	Mitigation strategy
Low quality of classification	Collect more data or introduce new features, including product pictures
Invalid product features or descriptions	Rule-based filters ensuring and enforcing reasonable inputs to the model and publicly tolerated contents of the product card
Change in category system	Design a correspondence between old and new categories. Retrain the model with new categories when sufficient sample can be collected.
Inappropriate descriptions designed to mislead the model, advertise some product, offend other customers, etc.	Rule based product card filtering
Other risks	
Risk	Mitigation strategy
Competitors deploy similar service	Introduce more features into analysis - add product photos into classification pipeline

Table 3: Main project risks and their mitigation strategy

2.6.4 Costs and benefits

Major expense categories include: project development and maintenance; periodic model adjustment for changes in data and/or company or governmental policies; monthly expenses: electricity, internet traffic.

As it is more a proof of concept project rather than a production ready system we expect to spend about 300000 rubles for salaries of two developers. Ideally, the team also needs a security engineer and preferable a software architect and that doubles the expenses. Maintenance of the project can be performed on demand by the forces of on-site teams. If the team spends on average 1 week in a month to maintain the service, it will cost 75000 rubles monthly. Model adjustment is unlikely to happen often but if we assume it happen in each quarter, it will cost about 300000 rubles on development itself and 50000 extra to organize extraction and anonymization of the data. Additional traffic and computational demands are not that high unless we use a modern NLP approaches thus we ignore them in our calculations.

To sum up, the project estimated cost in nearest two months is 300000 rubles.

It's annual maintenance will cost about 2.3 million rubles. Notice that the current estimates do not include the expenses on the internet traffic and hardware platforms into calculation as the service that is to be developed is not computationally demanding and can be deployed on a currently available hardware.

In 2023 the growth rate of the number of sellers has significantly dropped because according to the reports about 53% of sellers quit marketplaces after the first year of their business. Among the main reasons 27% of sellers mention low sales due to inefficiency of their product cards and significant time needed to fill in the card. Therefore, we believe that simplification of the product card creation process might significantly reduce the rate of sellers churn.

Let us estimate the annual profit from this system: in 2023 number of sellers on Ozon platform increased by 21% and reached 450000. About 18000 sellers are likely to leave the platform by the end of the year. If the service will manage to keep just 10% of those sellers on a marketplace, it will generate 4 million rubles of sales given the average order price in 2023 of 2188 rubles. Platform takes 10% as a fees so it will produce 400 thousands rubles extra. In such a simple model we expect the project to be loss-making. However, we understand that there are many effects we could not directly take into account or estimate: if 0,1% of sellers will make not 1 but 2 sales in a year it will generate 9.9 millions of income extra thus making the project profitable; our current sellers are likely to get more committed to our marketplace as they can clearly see that we also take their needs into account and further they might invite their friends to start their business on our platform or switch from another marketplace; our customers are more likely to find product they are looking for in a right category thus returning more frequently to our marketplace.

Overall, we can see that the project has a moderate risk. In the worst case we will lose about 10 million rubles in 3 years which negligibly small with respect to 10 billions of income of Ozon in 2023. If, however, the project deploys successfully in the most sceptical scenario we expect the income to increase by 0.1% annually. If some of the aforementioned factors will start to work we might expect this ratio to increase at least 10 times.

2.7 Project plan

The first stage of the project takes one week and includes data quality assessment, data cleaning and imputation if needed.

The second stage takes two weeks and includes basic data understanding, its limitations, engineering of new features. The largest risk on this phase is non-representative data sample. All of the sudden in this case we can not take any actions to mitigate that risk, and the only action we may take if it becomes true is to help the customer to organize the aggregation of the appropriate data sample. This risk also includes an outdated data which is different in properties from a current data. If we can not obtain a fresh or high quality dataset then the designed model might fail in production system, so we will have either organize A/B testing and see how model performs in the real environment, or invite a group of sellers to test new system.

During the third stage that lasts for 3 weeks we will try different models and

assess their quality. The models also include different pre-processing steps like TF-IDF or Word2Vec embeddings. To ensure the interpretability we may use classical statistical models like decision trees and forests, logistic regression with one-vs-rest strategy, and simple shallow MLP. We might need to introduce new features and thus roll back to one of previous stages.

Finally, in the last week of the project we aim to organize the project deliverables in a format ready for deployment and presentation to management of the business. It might happen that other marketplace delivers similar product and we will have to adjust model to ensure its higher quality and convenience for the end user.

References

- [1] Tinkoff eCommerce и Data Insight. Рынок нишевых маркетплейсов в России, 2023. URL <https://www.tadviser.ru/images/9/99/Tinkoff-ecommerce-and-data-insight-research-niche-marketplace-market-in-russia.pdf>.
- [2] Nirmal NK. Online retail dataset, 2021. URL <https://www.kaggle.com/datasets/nirmalnk/online-retail-dataset>. Accessed: 2024/06/18.
- [3] Ozon Sellers. Ozon sellers. create commodity card, 2023. URL <https://seller-edu.ozon.ru/work-with-goods/zagruzka-tovarov/creating-goods/sozdanie-tovarov-v-lk>.